



Kernel Runtime Security Instrumentation

KP Singh

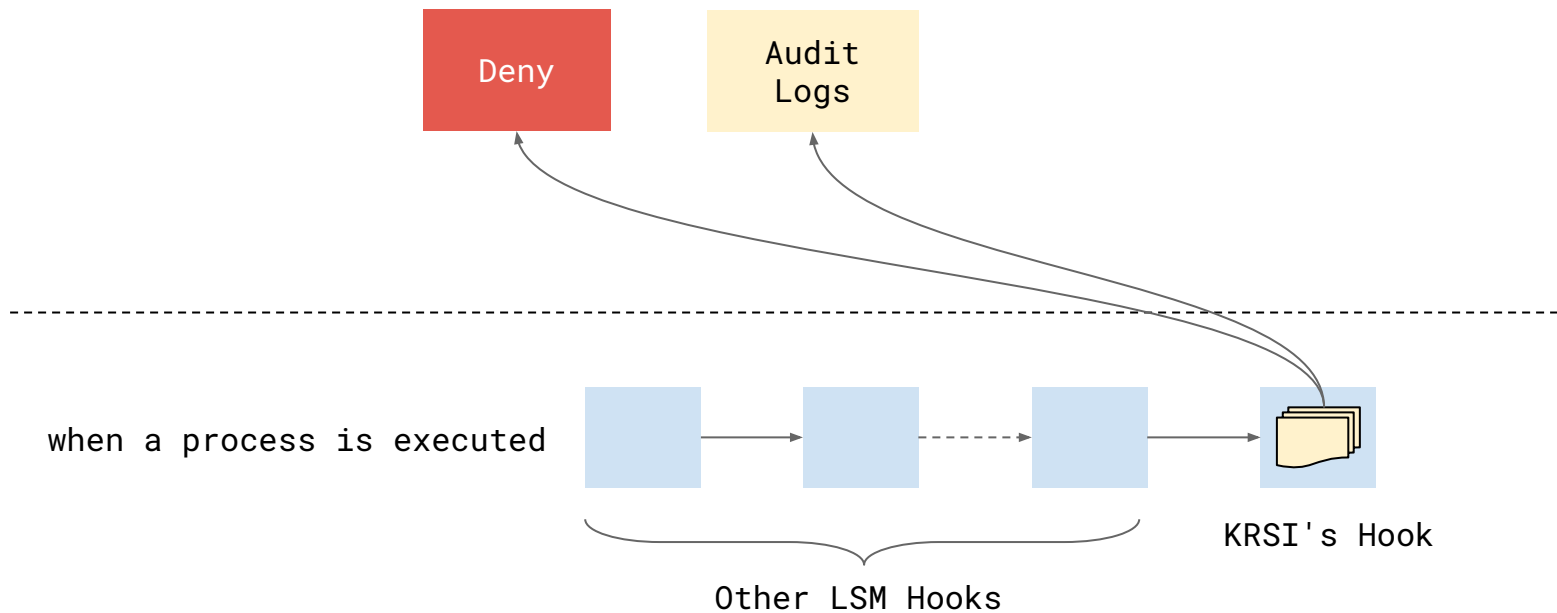
What?

 eBPF + LSM

=

KRSI

Audit + MAC





bytecode



bytecode

bpf()



 **BPF Verifier**

Approved



bytecode

bpf()



 **BPF Verifier**

Approved

 **BPF JIT**



x86_64



bytecode

bpf()

 BPF Verifier

Approved

 BPF JIT



x86_64

trampolines

BPF LSM Hook

BPF Trampolines: Types

fentry



```
int bpf_lsm_bprm_check_security  
{
```

fmod_ret



```
    <update ret by calling fmod_ret progs>
```

```
    if (ret != 0)  
        goto fexit;
```

```
original_function:
```


```
    <side_effects_happen_here>
```

fexit



```
}
```

BPF Trampolines: Types



LSM Hooks

LSM_HOOK Macros (lsm_hook_defs.h)

```
LSM_HOOK(  
    int,                // Return Type  
    0,                  // Default Return Value  
    bprm_check_security, // Name  
    struct linux_binprm *bprm // Parameters  
)
```

"Macro magic"

Default Callbacks

```
#define LSM_HOOK(RET, DEFAULT, NAME, ...) \
\
noinline RET bpf_lsm_##NAME(__VA_ARGS__) \
\
{ \
    return DEFAULT; \
}

#include <linux/lsm_hook_defs.h>

#undef LSM_HOOK
```

Default Callbacks

[...]

```
noinline int  
bpf_lsm_bprm_check_security(struct linux_binprm *bprm)  
{  
    return 0;  
}
```

[...]

Initialize LSM Hooks

```
#define LSM_HOOK(RET, DEFAULT, NAME, ...) \
    LSM_HOOK_INIT(NAME, bpf_lsm_##NAME),
#include linux/lsm_hook_defs.h>

#undef LSM_HOOK
```


Initialize BPF LSM Hooks

```
[...]
```

```
LSM_HOOK_INIT(bprm_check_security,  
              bpf_lsm_bprm_check_security);
```

```
[...]
```

Implementing Hooks

Load a program for `bprm_check_security`

BPF Program

```
SEC("lsm/bprm_check_security")
int BPF_PROG(my_prog, struct linux_binprm *bprm, int ret)
{
    __u32 pid = bpf_get_current_pid_tgid() >> 32;

    if (monitored_pid == pid)
        bprm_count++;

    return 0;
}
```

Context Simplification

`int *ctx`

`int ret`

`struct linux_binprm *bprm`

`BPF_PROG(my_prog, struct linux_binprm *bprm, int ret)`

```
my_prog(int *ctx) {  
    __my_prog(ctx[0], ctx[1])  
}
```

Verification

/sys/kernel/btf/vmlinux

Compact type information
(BTF)



~125MB of
DWARF



btfxctx_access

Verifier



BPF LSM Hooks: Object File

```
$ objdump -Sr kernel/bpf/bpf_lsm.o
```

```
LSM_HOOK(int, 0, bprm_check_security, struct linux_binprm *bprm)
```

```
100:    e8 00 00 00 00    callq 105 <bpf_lsm_bprm_check_security+0x5>
101:                                R_X86_64_PLT32 __fentry__-0x4
105:    31 c0            xor     %eax,%eax>:
107:    c3             retq
108:    0f 1f 84 00 00 00 00    nopl   0x0(%rax,%rax,1)
```

BPF LSM Hooks: after __init

```
LSM_HOOK(int, 0, bprm_check_security, struct linux_binprm *bprm)
```

```
100:  0f 1f 44 00 00      nopl 0x00(%eax,%eax,1)
```

```
105:  31 c0              xor    %eax,%eax>:
```

```
107:  c3                retq
```

```
108:  0f 1f 84 00 00 00 00  nopl 0x0(%rax,%rax,1)
```

ftrace_nop_initialize



BPF Trampoline Update

```
LSM_HOOK(int, 0, bprm_check_security, struct linux_binprm *bprm)
```

```
100:  e8 00 00 00 00 64      callq  <trampoline_image>
```

```
105:  31 c0                  xor     %eax,%eax:
```

```
107:  c3                    retq
```

```
108:  0f 1f 84 00 00 00 00  nopl    0x0(%rax,%rax,1)
```

```
200: <trampoline_image>
```



arch_prepare_bpf_trampoline

LSM Trampolines: Creation

```
push    %rbp
mov     %rsp,%rbp
sub     $0x10,%rsp
push    %rbx
```

Create a frame for a stack size of 16 (0x10) bytes:

- 8 bytes for `struct linux_binprm *bprm`
- 8 bytes to save the return value

LSM Trampolines: Invocation

```
mov    %rdi, -0x10(%rbp)
```

} Save the first argument on the stack

```
xor     %eax, %eax
```

```
mov     %rax, -0x8(%rbp)
```

} Clear out the return value passed to first LSM program.

```
callq   __bpf_prog_enter
```

```
mov     %rax, %rbx
```

```
lea     -0x10(%rbp), %rdi
```

} ctx (int * array) for the BPF program

```
callq   addr_of_jited_lsm_prog
```

```
mov     %rax, -0x8(%rbp)
```

} Call the JITed program
Save the return value on the stack

```
movabs  $addr_struct_bpf_prog, %rdi
```

```
mov     %rbx, %rsi
```

```
callq   __bpf_prog_exit
```

LSM Trampolines: BPF_MODIFY_RETURN

```
cmpq    $0x0, -0x8(%rbp)
jne     <do_exit>
```

} Skip calling the original function
and the rest of the programs upon
a non-zero return value

[...]

```
mov     -0x10(%rbp), %rdi
callq   <bpf_lsm_bprm_check_security+0x5>
mov     %rax, -0x8(%rbp)
```

```
nopl    0x0(%rax,%rax,1)
nopw    0x0(%rax,%rax,1)
```

} nops to align jump target

do_exit:

BPF Trampolines - Exit

```
mov    -0x8(%rbp),%rax
```

} Update the return value from the stack

```
pop     %rbx  
leaveq  
add     $0x8,%rsp  
retq
```

Improvements

Indirect Calls

```
hlist_for_each_entry(P, &security_hook_heads.FUNC, list) \
{\
    RC = P->hook.FUNC(__VA_ARGS__); \
    if (RC != 0) \
        break; \
}
```



Indirect calls worsened by
retpolines

and default callbacks are added
everywhere!!

So How bad is it?

```
int main(void) {  
    int fd = eventfd(0, 0);  
    int c = 10000;  
  
    while (c--)  
        eventfd_write(fd, 1);  
  
    return 0;  
}
```

```
int main(void) {  
    int fd = eventfd(0, 0);  
    int c = 10000;  
  
    while (c--)  
        eventfd_write(fd, 1);  
  
    return 0;  
}
```



```
int main(void) {  
    int fd = eventfd(0, 0);  
    int c = 10000;
```

```
    while (c--)  
        eventfd_write(fd, 1);
```



+ ~4%

```
    return 0;
```

```
}
```

We know the addresses of LSM Hooks
at `__init..`



Use DEFINE_STATIC_CALL...



Slots for call instructions at
compile time.

SELINUX_SLOT

nop

AA_SLOT

nop

•

•

BPF_SLOT

nop

bprm_check_security call slots patched
at __init

SELINUX_SLOT

call selinux_bprm_check_security

AA_SLOT

call aa_bprm_check_security

•

•

BPF_SLOT

nop

Patched only
when there is a
BPF program
attached.

Progress on DEFINE_STATIC_CALL
is slow...



Upcoming..

BPF Ring Buffer

Merged!

Storage blobs a.k.a
`bpf_local_storage`

patches on the list!

Sleepable BPF

patches on the list!

bpf_d_path helper

patches on the list!

Advanced string helpers
(argv, file paths..)

Not started
(Custom Patches)

Load BPF programs during boot..

Not started...

Thank You!